M7 Practical Challenge: Predictive Modeling with DataRobot

Table of Contents

Alejandro C. Parra Garcia

Input Data

The data is the same used in the previous task: Cars.csv Containing information about cars and its prices.

Once imported to Data Robot the data is presented as follow:

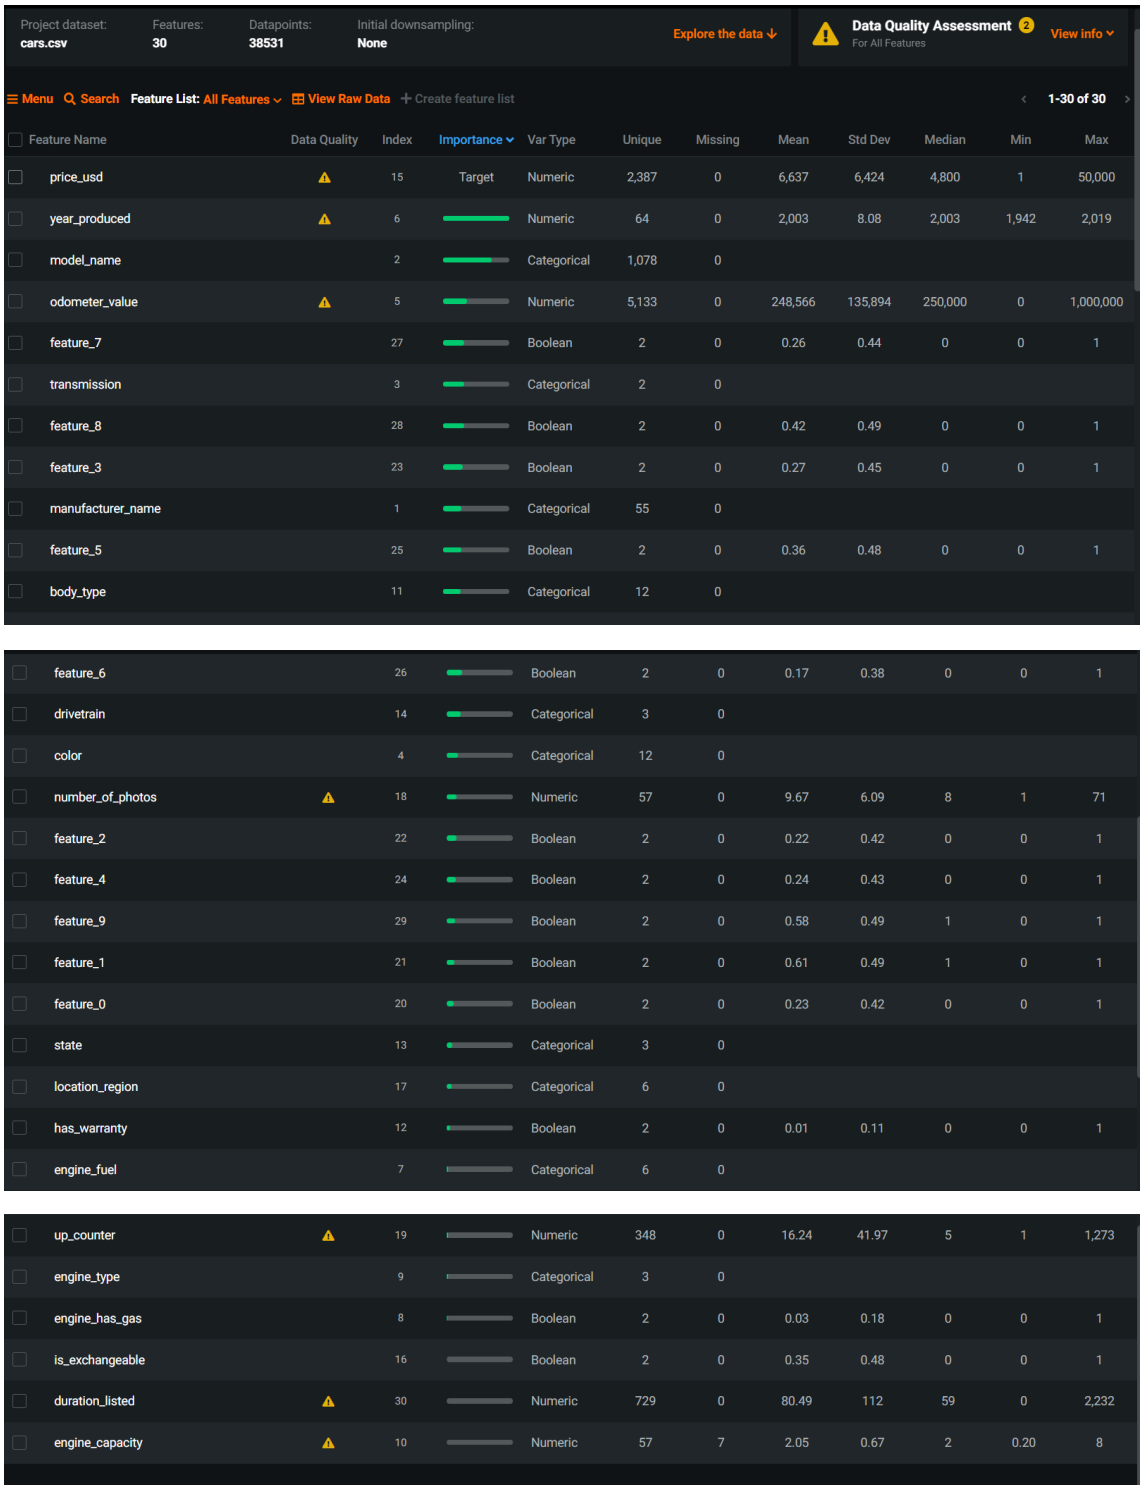

Project dataset: cars.csv	Features: 30	Datapoints: 38531	Initial downsampling: None	Explore the data	Data Quality Assessment 2 For All Features	View info

Menu | Search | Feature List: All Features | View Raw Data | Create feature list | 1-30 of 30

Feature Name	Data Quality	Index	Importance	Var Type	Unique	Missing	Mean	Std Dev	Median	Min	Max
price_usd	⚠	15	Target	Numeric	2,387	0	6,637	6,424	4,800	1	50,000
year_produced	⚠	6		Numeric	64	0	2,003	8.08	2,003	1,942	2,019
model_name		2		Categorical	1,078	0					
odometer_value	⚠	5		Numeric	5,133	0	248,566	135,894	250,000	0	1,000,000
feature_7		27		Boolean	2	0	0.26	0.44	0	0	1
transmission		3		Categorical	2	0					
feature_8		28		Boolean	2	0	0.42	0.49	0	0	1
feature_3		23		Boolean	2	0	0.27	0.45	0	0	1
manufacturer_name		1		Categorical	55	0					
feature_5		25		Boolean	2	0	0.36	0.48	0	0	1
body_type		11		Categorical	12	0					
feature_6		26		Boolean	2	0	0.17	0.38	0	0	1
drivetrain		14		Categorical	3	0					
color		4		Categorical	12	0					
number_of_photos	⚠	18		Numeric	57	0	9.67	6.09	8	1	71
feature_2		22		Boolean	2	0	0.22	0.42	0	0	1
feature_4		24		Boolean	2	0	0.24	0.43	0	0	1
feature_9		29		Boolean	2	0	0.58	0.49	1	0	1
feature_1		21		Boolean	2	0	0.61	0.49	1	0	1
feature_0		20		Boolean	2	0	0.23	0.42	0	0	1
state		13		Categorical	3	0					
location_region		17		Categorical	6	0					
has_warranty		12		Boolean	2	0	0.01	0.11	0	0	1
engine_fuel		7		Categorical	6	0					
up_counter	⚠	19		Numeric	348	0	16.24	41.97	5	1	1,273
engine_type		9		Categorical	3	0					
engine_has_gas		8		Boolean	2	0	0.03	0.18	0	0	1
is_exchangeable		16		Boolean	2	0	0.35	0.48	0	0	1
duration_listed	⚠	30		Numeric	729	0	80.49	112	59	0	2,232
engine_capacity	⚠	10		Numeric	57	7	2.05	0.67	2	0.20	8

Then the Target Variable is selected as price_usd:

Model Selection

As we can see the best model is eXtreme Gradient Boosted Trees Regressor with Early Stopping (Poisson Loss), with an RMSE of 1746.

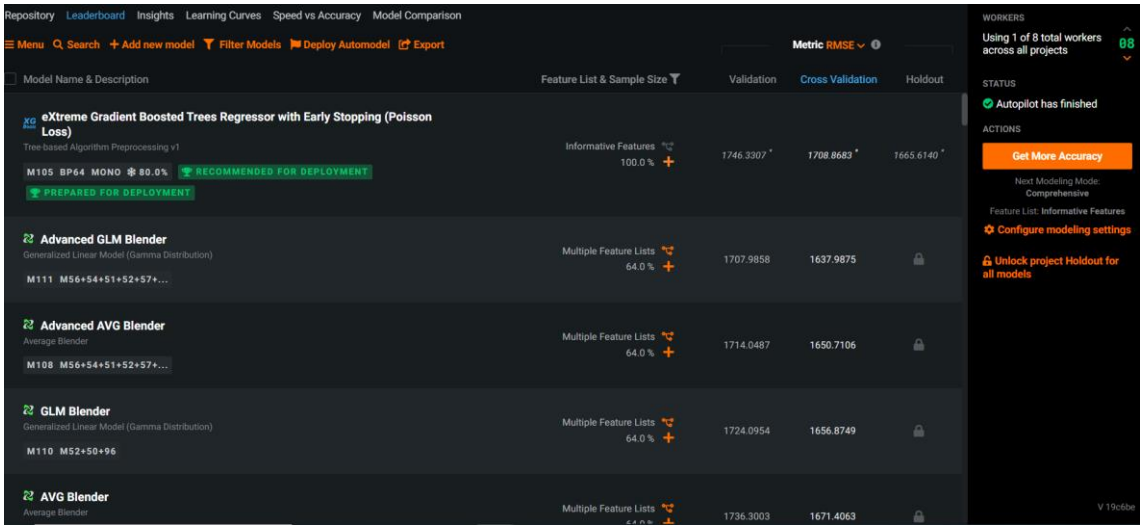

eXtreme Gradient Boosted Trees Regressor with Early Stopping (Poisson Loss)

Describe Section

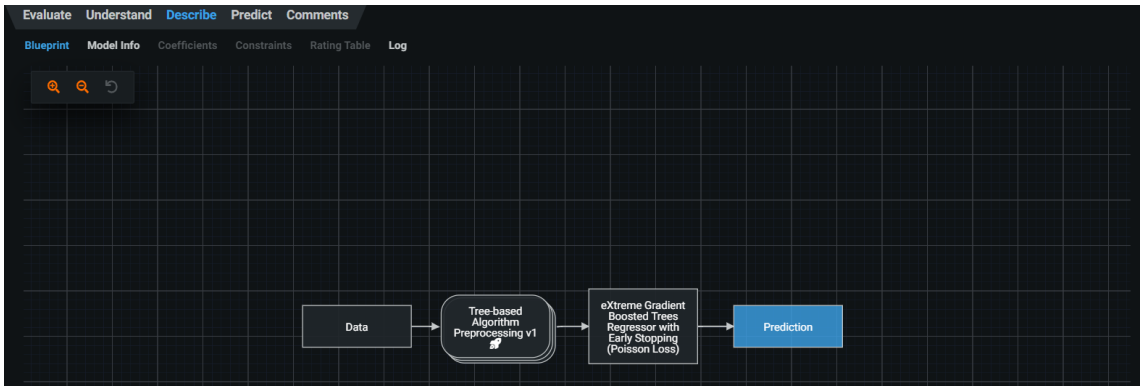

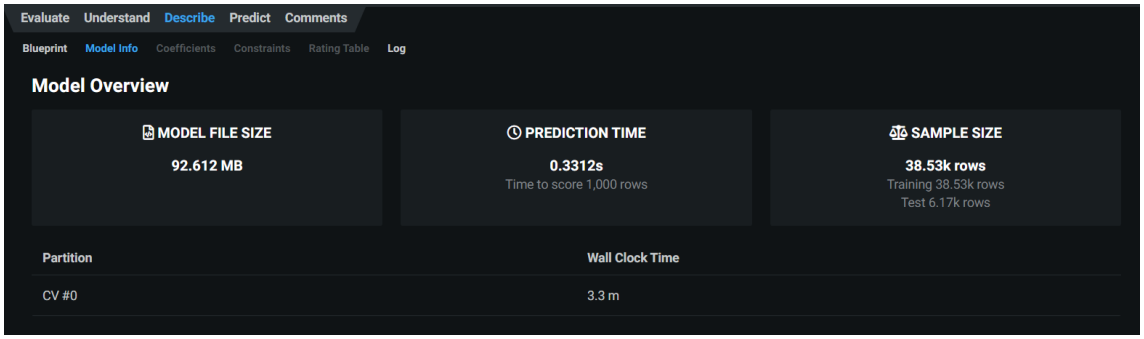

As we can see, the model is able to run 1000 rows of data in 0.33 seconds. Meaning it can run a 1 million data sample in less than 6 minutes. Making it a really fast model.

Evaluate Section

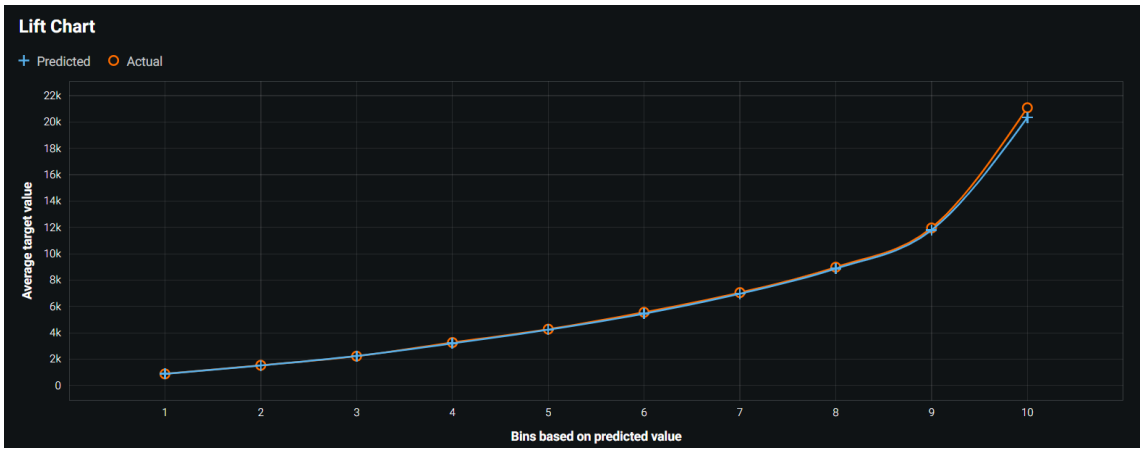

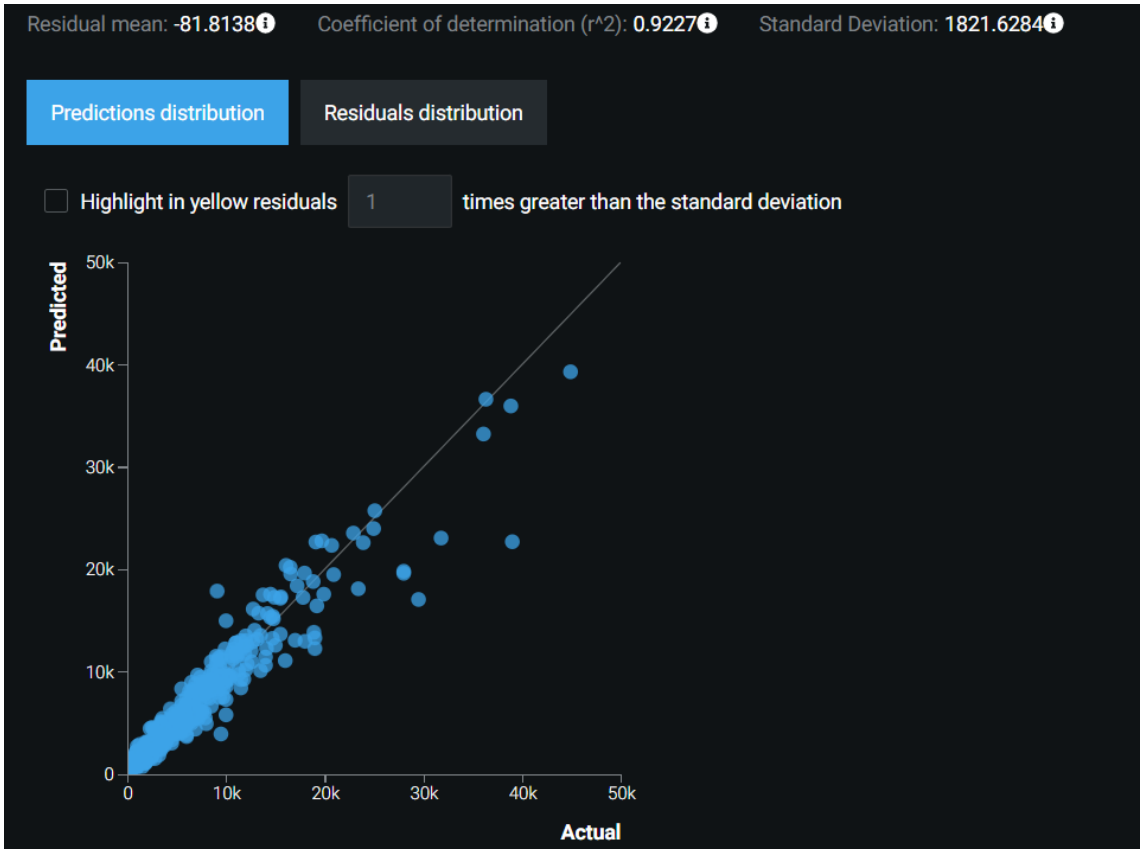

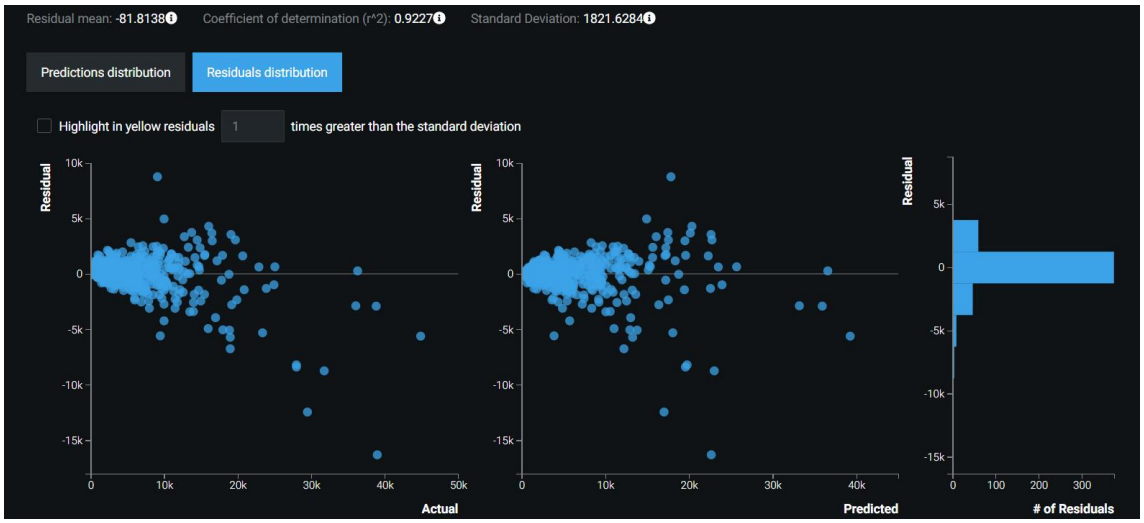

As we can see with the graphs is that the model is more precise with the lower prices' cars. But this can be because there are less datapoints for the higher prices' cars.

Feature Fit

We can see the importance of each variable on the target variable:

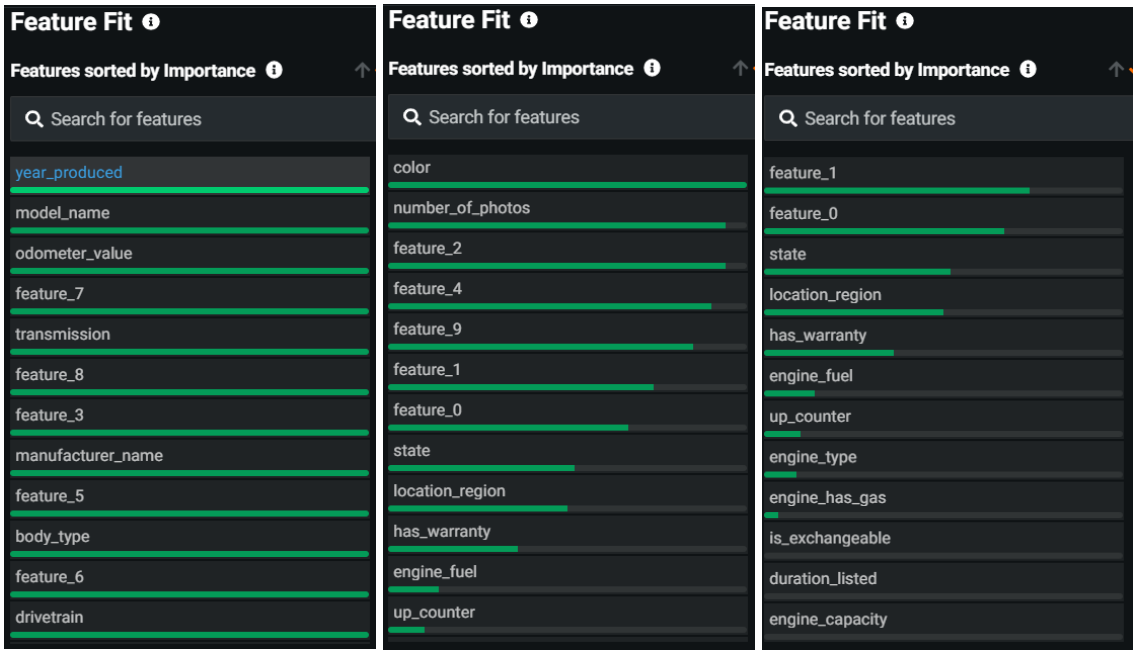

We can see that there are some variables that are extremely important, such as: year_produced, model_name, odometer_value, feature_7, transmission, feature_8, feature_3, manufacturer_name, feature_5, body_type, feature_6, drivetrain, color. Then the importance of each variable diminishes, until the last variables that have no impact on the model. We can see a couple of graphs: (year_produced, odometer_value, color)

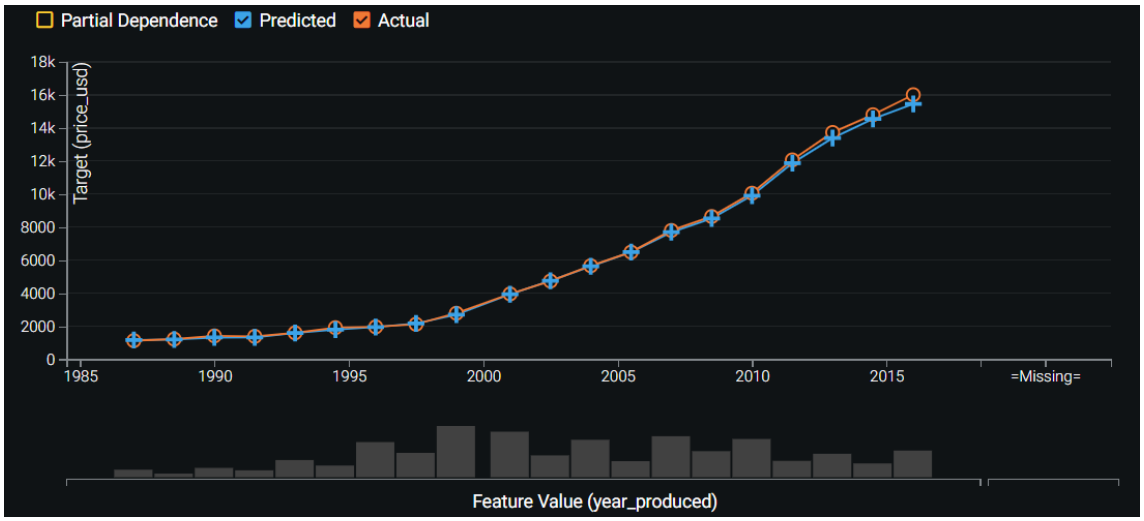

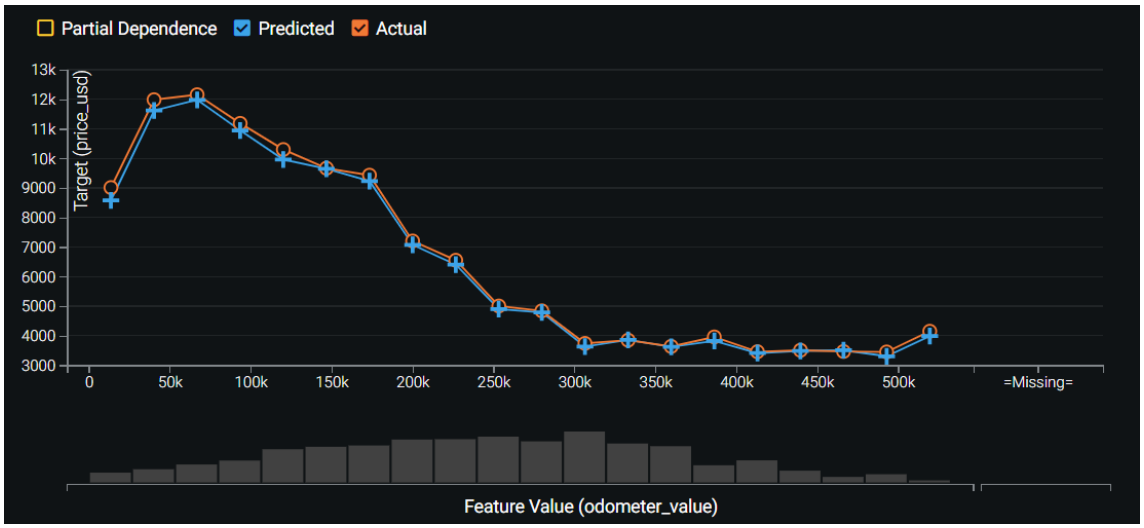

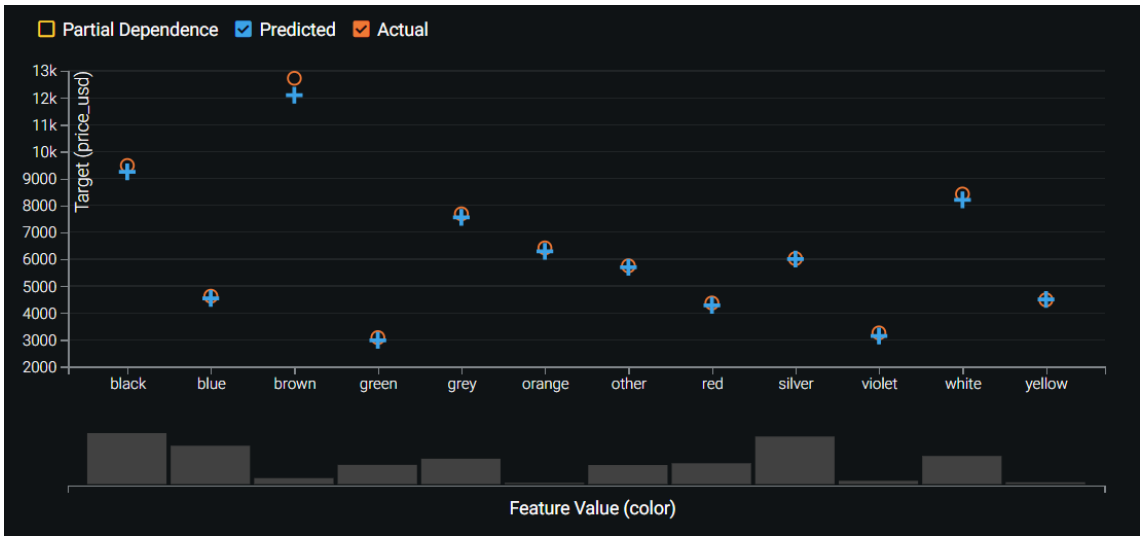

We can see, that the older the car is, the lower its price is, we can also see that the number of kilometres travel by a car have an invers relation to the price, the higher the number, the lower the price. The color also have an impact on the value, as shown on the graph, different color having a different impact on the price.

Understand section

Feature Impact

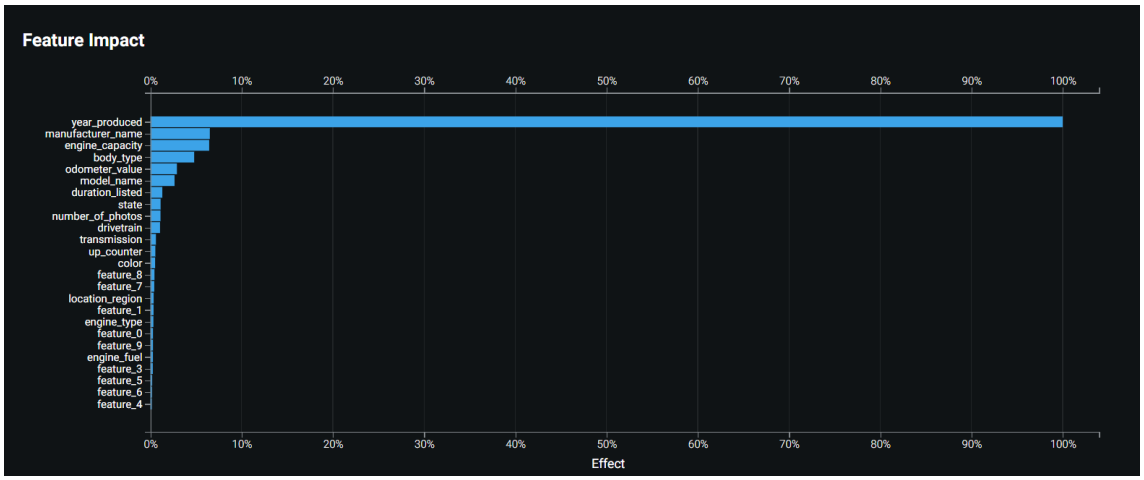

We can also check the importance of each feature to the model. We can see that the year have the highest importance, similar to what we mention on the previous sections.

Hypertuning

So, now we are going to change some parameters to try to improve the performance:

Learning rate

We are trying: 0.01, 0.1, 0.2

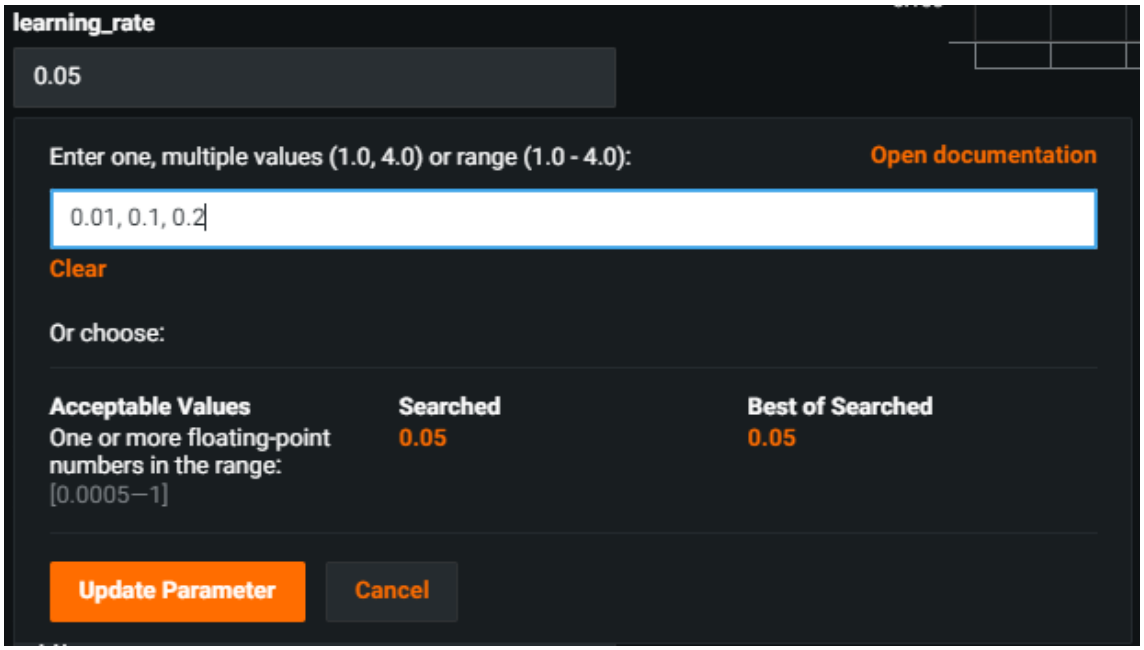

Max depth

We are trying: 3, 5, 9, 11

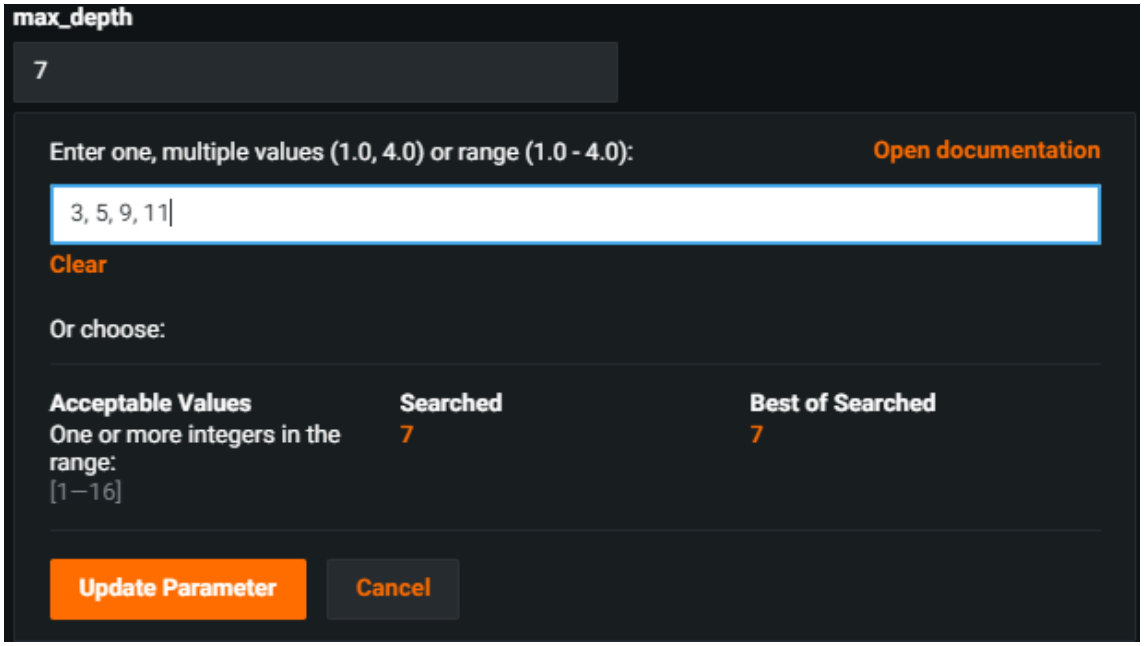

Once we test with this parameter, we get the next results:

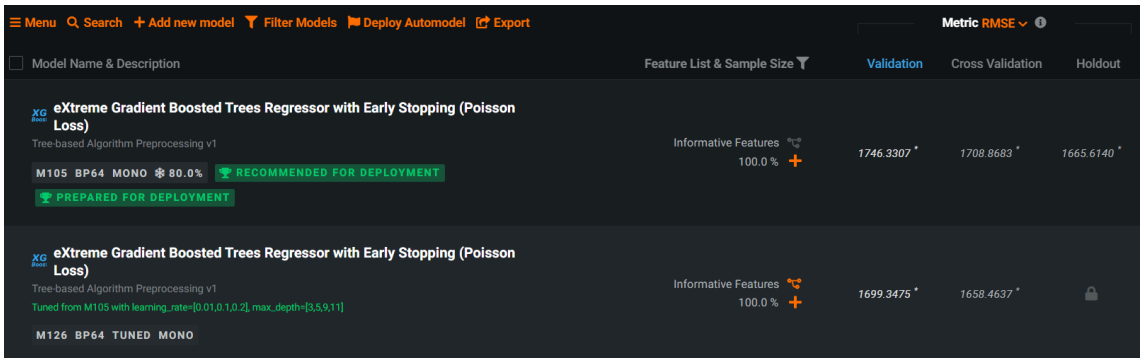

We can see that the new model achieved a better Validation and Cross Validation metrics. So now we need to check which values get the best results, for that we now go to check the model itself and check the values. We get the next results.

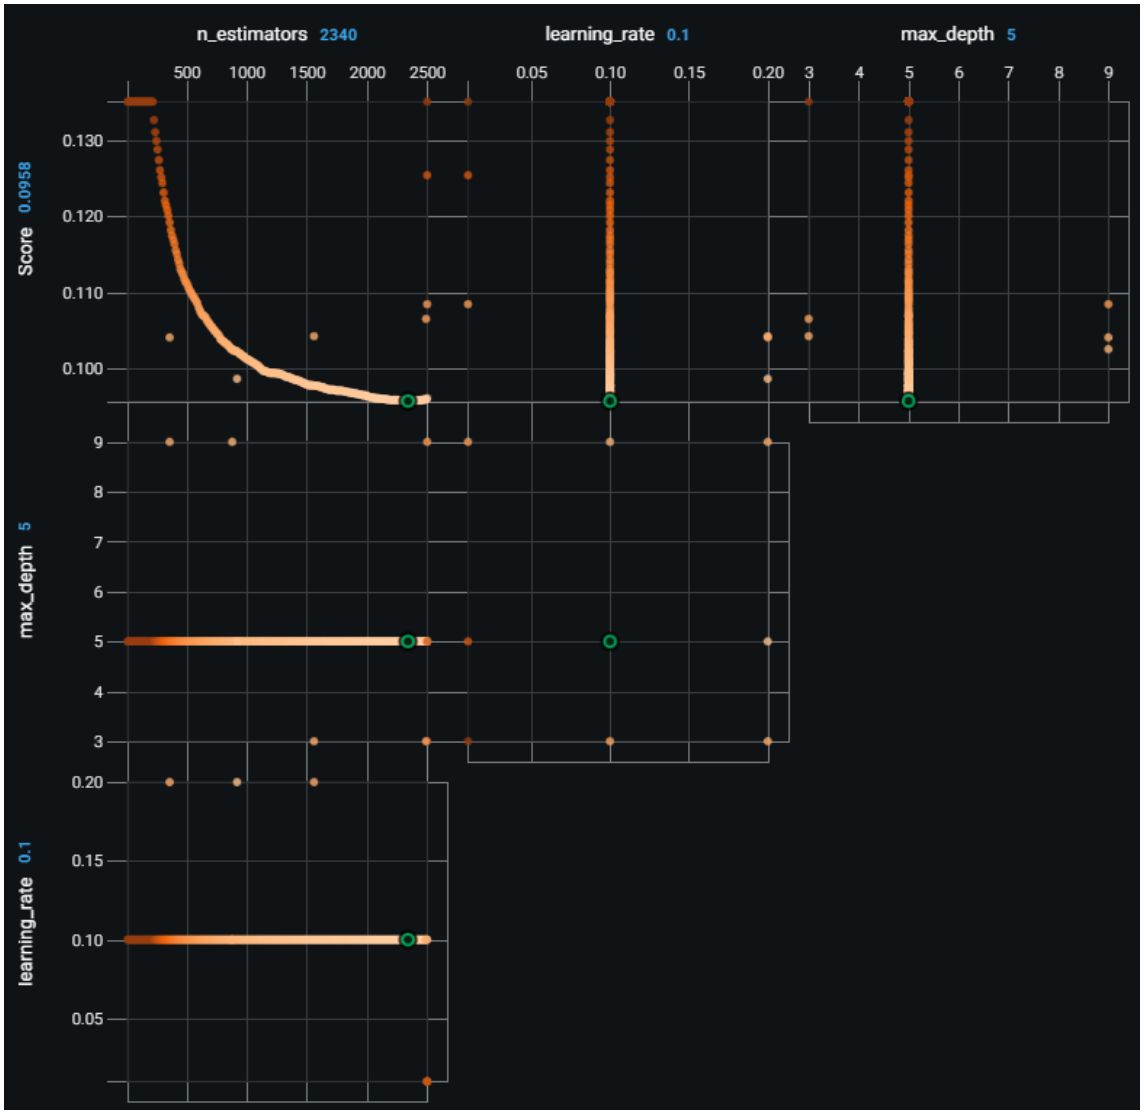

We see that the best performance is achieved with a learning rate of 0.1 and a max_depth of 5. With this metrics the RMSE of the cross validation of 1658.46

If we now check the lift chart of the new model, as shown on the next graph, and compare it with the previous one, we see that the model achive a higher precision in the most expensive cars.

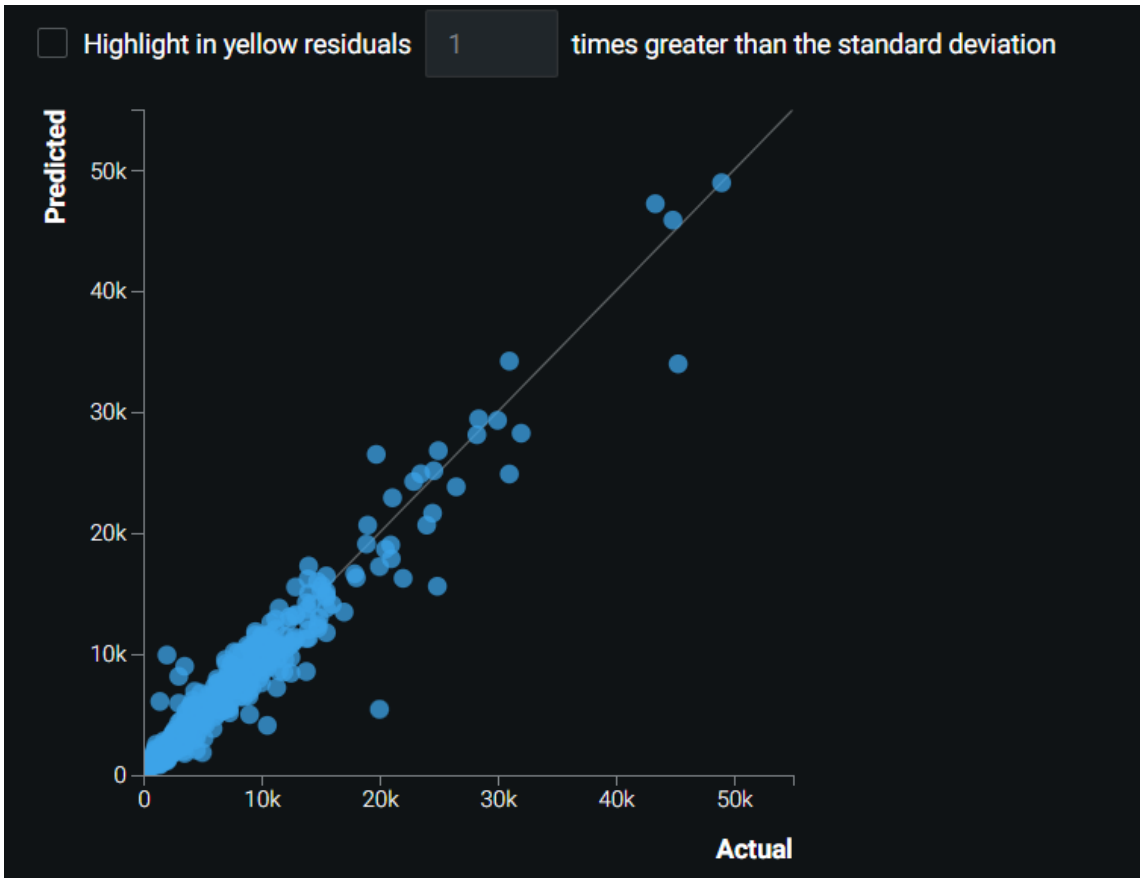

Benefits/Concerns

Some of the Benefits, are the extremely easy nature of the program, since you need to only drag and drop to get the results, it is also extremely fast, as it can process multiple different models to get the highest performance. Combining these two features you get a powerful tool, easy to used and fast for creating predictive models.

The concerns revolve around the fact that the software is employed online, meaning that in the event of failure, you lose connections to your built models. Also, in training the data, you need to uploaded, which can be risky or illegal in the case of sensitive information.